

Introduction to Running Computations on the High Performance Clusters at the Center for Computational Research

Cynthia Cornelius

Center for Computational Research
University at Buffalo, SUNY

`cdc at buffalo.edu`

Spring 2015

CCR Resources

The Center for Computational Research provides high performance computing resources to the University at Buffalo.

- Supporting faculty research and classroom education, as well as local business and University collaborations.
- High performance and high through-put cluster.
- High performance remote visualization.

CCR Cluster

- Users run programs on the cluster as **jobs** submitted to queues.
 - A **job** is a request for computers, cores, memory, networking, queue, and time.
- The CCR cluster is a high performance.
 - Large parallel jobs run on the cluster.
- The CCR cluster is high through-put.
 - A large number of jobs run on the cluster at one time.

CCR Cluster

- The CCR cluster is collection of linux computers, private networks, a shared file system, a login computer, and job scheduler.
- Resources must be requested from the **job scheduler**.
- The user must decide what resources are required for the job.
- The user must control and optimize the job.

CCR Cluster

- The CCR cluster is **NOT** a **cloud**.
- Resources are NOT on demand.
- There is NO rapid elasticity.
- There is NO measured service to optimize the use of resources.

CCR Cluster

- The CCR cluster provides over 8,000 cores.
- There is a mix of 8-core, 12-core , 16-core and 32-core compute nodes.
 - A **compute node** is a linux machine with memory, disks, and cores (cpus), as well as both ethernet and Infiniband network connections.
- The GPFS (IBM General Parallel File System) is used for storage.
- There is a login front-end server.
- SLURM (Simple Linux Utility Resource Manager) is the job scheduler.

CCR Cluster

- **IMPORTANT!**
- Recent and planned changes to the CCR cluster.
 - All users should look at the [2015 What's New](#) guide.

Where are my files?

- GPFS provides storage space of user home directories, projects directories and global scratch space.
- All compute node access GPFS.
- A home directory has a 2GB quota by default. The path is **/user/<username>**.
- Faculty can request projects space for the research group. The path is **/projects/<faculty-username>**
- Home and project directories are backed up daily.

Where are my files?

- **/gpfs/scratch** is global scratch space.
 - This space is for temporary use.
 - Use scratch space for running jobs only.
 - Files older than 3 weeks are automatically removed.
 - There is no backup of files in **/gpfs/scratch**.
- Every compute node has a **/scratch** directory on the local disk. Jobs can use this local scratch directory.

CCR Cluster Compute Nodes

- There are over 700 compute nodes providing ~8000 cores.
- Compute nodes are grouped according to number of cores.
- A job will always be assigned nodes with the same number of cores.
 - No job would ever have a mix of 8-core and 12-core compute nodes.
- The maximum number of compute nodes that users can request in a job is the total for that group of nodes.

Number of Compute Nodes

- **12-core** compute nodes — **372**
- **8-core** compute nodes — **256**
- **16-core** compute nodes — **32**
- **32-core** compute nodes — **18**
- **12-core** nodes with GPUs — **32**
- [more on compute nodes](#)

How to choose a compute node

- Most users choose compute nodes based on number of nodes, number of cores, and memory per node required for the job.
- 8-core compute node have 24GB of memory.
- 12-core compute nodes have 48GB of memory.
- 16-core compute nodes have 128GB of memory.
- 16 of the 32-core compute nodes have 256GB, while 2 have 512GB of memory.
- GPU Compute nodes have 48GB of memory and 2 Nvidia Fermi GPUs.

How to choose a partition

- SLURM **job queues** are referred to as partitions. Here are the **partitions** for the CCR cluster.
- **general-compute** - default partition if no partition is specified for a job.
 - almost all compute nodes
 - **per user limit of 1000** running and pending jobs at a time
 - maximum time limit of 72 hours

How to choose a partition

- **gpu** - higher priority for only the compute nodes with GPUs.
 - **per user limit of 32** running or pending jobs at a time.
 - maximum time limit of 72 hours
- **largemem** - higher priority for only the 32-core compute nodes.
 - **per user limit of 32** running or pending jobs at a time.
 - maximum time limit of 72 hours

How to choose a partition

- **debug** - small partition for quick debugging.
 - 8 dedicated compute nodes
 - 4 8-core compute nodes,
 - 2 12-core compute nodes
 - 1 16-core node with 2 GPUs
 - 1 16-core node with a XEON PHI coprocessor
 - **per user limit of 4** running or pending jobs at a time.
 - maximum time limit of **1 hour**

How to choose a partition

- **viz** - partition for jobs running on the remote visualization nodes.
 - Users do not submit directly to the viz partition.
 - **per user limit of 32** running or pending jobs at a time.
 - maximum time limit of 24 hours
 - [how to use the Remote Visualization compute nodes](#)

How to choose a partition

- **supporters** - higher priority partition for users belonging to research groups that have contributed funds to CCR.
 - almost all compute nodes
 - **per user limit of 1000** running or pending jobs at a time
 - maximum time limit of 72 hours
- [more on CCR SLURM partitions](#)

Access to the CCR Cluster

- The front-end machine is **rush.ccr.buffalo.edu**
- 32-core node with 256GB of memory.
- Accessible from **UB network only**.
- Only secure protocols, such as **ssh** and **sftp**, are permitted to access the front-end.
- Small test computations can run on the front-end machine.
 - CPU time limit of 15 minutes on the front-end. Processes that exceed the time limit are automatically terminated.

How to login

- Login from Linux or Mac
- `ssh -X rush.ccr.buffalo.edu`
- **`ssh -X UBITusername@rush.ccr.buffalo.edu`**
 - The -X flag enables a graphical display. This is optional.
- Windows users must install X-Win32 or PuTTY for the secure login.
- The **X-Win32** program allows for a graphical display.
 - [more on how to login](#)

How to transfer files

- **Filezilla** is graphical file transfer program, which is available for Windows, Linux and MAC computers.
- **UBVPN** must be used to access the rush login machine from off campus.
- Users should setup UB-Secure on their laptops for connecting through wireless.
- [more on how to transfer files](#)
- Get the software from the UBit webpage.
 - [UBit software](#)

Command Line Environment

- The compute nodes and front-end machine run Linux. It is a command line **UNIX environment**.
- Users should have know basic UNIX commands, such as ls, cd and mkdir.
- There are several editors available: emacs, nano and vi.
- Use the dos2unix command to remove any hidden characters in text files transferred from Windows machines.
- [more on basic UNIX commands](#)
- [CCR UNIX Reference Card](#)

Using Software and Compilers

- Compilers, JAVA, Python, MPI and application software are available to use on cluster.
- [Lmod](#) is used to set paths and variables for the software.
- “**module avail**” shows all installed software packages.
- “**module load <package-name>**” puts the software in the user’s path.
- [more on using Lmod and modules](#)

Compilers, MPI and more

- The Intel, PGI and GNU compilers are available on the cluster. GNU compilers are already in the user's default path.
- MPI (Message Passing Interface) implementations are installed on the cluster.
- More recent versions of JAVA are available.
- Anaconda Python is available.
 - [more on using Anaconda Python](#)

SLURM Commands

- **squeue** — shows the status of jobs.
- **sbatch** — submits a script job.
- **scancel** — cancels a running or pending job.
- **sinfo** — provides information on partitions and nodes.
- **snodes** — shows details of the compute nodes.
- **slurmjobvis** – graphical job monitoring tool.
- [more on SLURM commands](#)

How to submit an interactive job

- Submit an interactive job using the **fishbatch** wrapper.
- Useful for debugging.
- Specify partition, nodes, cores or tasks, and time.
- Once the job starts the user is logged into the compute node.
- [more on submitting an interactive job](#)

SLURM script example

```
#!/bin/sh  
#SBATCH --partition=general-compute  
#SBATCH --time=00:15:00  
#SBATCH --nodes=2  
#SBATCH --ntasks-per-node=8  
#SBATCH --mem=24000  
# Memory per node specification is in  
MB. It is optional.  
# The default limit is 3000MB per core.
```

SLURM script example

```
#SBATCH --job-name="hello_test"
#SBATCH --output=test-srun.out
#SBATCH --mail-
user=username@buffalo.edu
#SBATCH --mail-type=ALL
echo "SLURM_JOBID="$SLURM_JOBID
echo "SLURM_JOB_NODELIST"=
$SLURM_JOB_NODELIST
```

SLURM script example

```
echo "SLURM_NNODES"=$SLURM_NNODES
```

```
echo "SLURMTMPDIR"=$SLURMTMPDIR
```

```
echo "working directory =
```

```
"$SLURM_SUBMIT_DIR
```

```
module load intel/13.1
```

```
module load intel-mpi/4.1.3
```

```
module list
```

```
ulimit -s unlimited
```

SLURM script example

```
NPROCS=`srun --nodes=${SLURM_NNODES}  
bash -c 'hostname' |wc -l`  
echo NPROCS=$NPROCS  
echo "Launch helloworld with srun"  
#The PMI library is necessary for srun  
export I_MPI_PMI_LIBRARY=/usr/lib64/  
libpmi.so  
srun ./helloworld  
#  
echo "All Done!"
```

Commonly used SLURM variables

- **\$SLURM_JOBID**
- **\$SLURM_JOB_NODELIST**
 - Node list in SLURM format; for example f16n[04,06].
- **\$SLURM_NNODES**
 - Number of nodes
- **\$SLURMTMPDIR**
 - /scratch/jobid
 - local to the compute node
- **\$SLURM_SUBMIT_DIR**
 - Directory from which the job was submitted
- **NOTE!** Jobs start in the \$SLURM_SUBMIT_DIR.

Task Launching

- The number of cores/processes can be computed.
 - Use `srun` to get the total number of cores.
 - `NPROCS=`srun --nodes=${SLURM_NNODES} bash -c 'hostname' |wc -l``
- Intel-MPI `mpirun` and `mpiexec` are SLURM aware.
- `srun` will execute a command across nodes.
 - Typically, this is the best choice for launching a parallel computation.

How to submit a SLURM script job

- Submit an interactive job using the **sbatch** <your_slurm_script>
- The job will be submitted to the SLURM scheduler.
- The job will wait in the queue until the scheduler assigns resources to it. This is a pending state.
- [more on submitting a job script](#)

Check the status of the job

queue -u <username>

queue -j <job_id>

JOBID PARTITION NAME USER ST TIME NODES
NODELIST(REASON)

4832 general-c hello_te cdc R 0:20 2 f16n[10-11]

- Job status:
 - **R** – job is running.
 - **PD** – job is waiting for resource.
 - Reasons are usually (Resources) or (Priority).
 - Others commons reasons are CA (cancelled) and CD (completed).

Partition and node status

sinfo -p general-compute

```
PARTITION      AVAIL TIMELIMIT NODES  STATE
NODELIST
general-comput* up 3-00:00:00  264  idle
d07n07s[01-02],d07n08s[01-02], ...
```

- **Node states:**
 - **alloc** – all cores are in use.
 - **mix** – some cores are available.
 - **idle** – node is free. All cores are available.
 - **down**- node is down.
 - **drained** – node is offline.

Node status and details

```
[cdc@rush:~]$ snodes k14n16s01
```

HOSTNAMES	STATE	CPUS	S:C:T	CPUS(A/I/O/T)	CPU_LOAD
MEMORY	GRES	PARTITION	FEATURES		
k14n16s01	alloc	12	2:6:1	12/0/0/12	12.06 48000
(null)	general-compute*	IB,CPU-E5645			

```
cdc@rush:~]$ snodes k05n22
```

HOSTNAMES	STATE	CPUS	S:C:T	CPUS(A/I/O/T)	CPU_LOAD
MEMORY	GRES	PARTITION	FEATURES		
k05n22	idle	16	2:8:1	0/16/0/16	0.04 128000 mic:1
debug	CPU-E5-2660,MIC				

```
[cdc@rush:~]$
```

Node Sharing


- Compute nodes are **shared** among different jobs and users.
- Tasks are limited to the number of cores and memory specified.
- The integration of CPUSSETS and SLURM makes this possible.
 - **CPUSET** is a Linux kernel level mechanism that can be used to control access to individual cores.
- The default memory limit per core is 3000MB.

Node Sharing

- `--mem=24000`
 - Requests 24GB per node.
- `--mem-per-core=16000`
 - Requests 16GB on a core; use for serial job.
- Jobs exceeding the memory limit will be terminated by the resource manager.
- Check the [General Compute Cluster](#) webpage for node memory and core details.
- The `--exclusive` flag will request the nodes as dedicated. The nodes will not be shared.

Job monitoring

- The **slurmjobvis** is a graphical display of the activity on the node.
- CPU, memory, network, as well as GPU utilization, are displayed.



Putting it all together

- Login to the cluster login front-end machine.
- Transfer files to home directory.
- Locate necessary software with “module avail”
- Load modules and compile the code.
- Create a job script.
- Submit a script.
- Monitor job with `squeue` and `slurmjobvis`.

More Information and Help

- [CCR SLURM](#) web pages
- **/gpfs/courses/class-notes** for sample jobs.
- More sample SLURM scripts can be found in the /util/academic/slurm-scripts directory on rush.
- [Compute Cluster](#) web page
- [Remote Visualization](#) web page
- Users can get assistance by sending an email to ccr-help@ccr.buffalo.edu.