

# CSE 487/587

## Data Intensive Computing

### Lecture 8: Pig-Latin

Vipin Chaudhary

[vipin@buffalo.edu](mailto:vipin@buffalo.edu)

**716.645.4740**  
**305 Davis Hall**

# Overview of Today's Lecture

- Pig Latin (lot of stuff borrowed from Yahoo Research!)
- Midterm
- Review

# Problems with Mapreduce?

- Restricted programming model
  - Only two (or three) phases
  - Job chain for long data flow
- Too many lines of code even for simple logic
  - How many lines do you have for word count?
  - Programmers are responsible for this

# Pig to the Rescue

- What is Pig?
  - open-source high-level dataflow system
  - Provides a simple language for queries and data manipulation,
  - Compiled into map-reduce jobs that are run on Hadoop
  - Combines the high-level data manipulation constructs of SQL with the procedural programming of map-reduce
    - Systems like Netezza, Teradata, Oracle Cluster DB are \$\$\$\$\$
- Why is it important?
  - Social media and cloud companies like Yahoo, Google and Microsoft are collecting enormous data
  - Some form of ad-hoc processing and analysis of all of this information is required

# Where does Pig Fit?

## RDBMS

- Parallel DB expensive
- Rigid schemas
- Sql query construction

## Best of SQL and Map-Reduce

**COMBINE**      high-level declarative querying with  
low-level procedural programming

## Mapreduce

- Custom code for common tasks
- Complex workarounds for many tasks that do not fit map->combine->reduce

# Language Features

- Several options for user-interaction
  - Interactive mode (console)
  - Batch mode (prepared script files containing Pig Latin commands)
  - Embedded mode (execute Pig Latin commands within a Java program)
- Built primarily for scan-centric workloads and read-only data analysis
  - Easily operates on both structured and schema-less, unstructured data
- Extensive UDF support
  - Available in Java, Python and Javascript currently
  - Can be written for filtering, grouping, per-tuple processing, loading and storing
  - Use Piggy Bank – repository of Java UDFs written by other users.

# Data Model

- Supports four basic types
  - Atom: a simple atomic value (*int, long, double, string*)
    - ex: 'John'
  - Tuple: a sequence of fields that can be any of the data types
    - ex: ('John', 21)
  - Bag: a collection of tuples of potentially varying structures, can contain duplicates
    - ex: {('John'), ('Mike', (21, 25))}
  - Map: an associative array, the key must be a *chararray* but the value can be any type

# Pig Latin Example

Table urls: (url,category, pagerank)

Find for each sufficiently large category, the average pagerank of high-pagerank urls in that category

## SQL:

```
SELECT category, AVG(pagerank)
FROM urls WHERE pagerank > 0.2
GROUP BY category HAVING COUNT(*) > 106
```

## Pig Latin:

```
good_urls = FILTER urls BY pagerank > 0.2;
groups = GROUP good_urls BY category;
big_groups = FILTER groups BY COUNT(good_urls)>106;
output = FOREACH big_groups GENERATE
        category, AVG(good_urls.pagerank);
```

## Give me top 10 URLs

```
groups = GROUP urls BY category
Output = FOREACH groups GENERATE
        category, top10(urls)
```



# Pig Latin Example

**Set of URLs of pages classifieds as spam but have high pagerank score.**

```
Spam_urls = FILTER urls BY isSpam(url)
Culprit_urls = FILTER spam_urls BY pagerank > 0.8
```

**Give me top 10 URLs**

```
groups = GROUP urls BY category
Output = FOREACH groups GENERATE
         category, top10(urls)
```

# Data Model

- Atom - simple atomic value (ie: number or string)
- Tuple
- Bag
- Map

`'alice'`

# Data Model

- Atom - simple atomic value (ie: number or string)
- Tuple – sequence of fields; each field any type
- Bag
- Map

`('alice', 'lakers')`

# Data Model

- Atom - simple atomic value (ie: number or string)
- Tuple – sequence of fields; each field any type
- Bag – Collection of tuples
  - Duplicates possible
  - Tuples in a bag can have different field lengths and field types
- Map

$$\left\{ \begin{array}{l} ('alice', 'lakers') \\ ('alice', ('iPod', 'apple')) \end{array} \right\}$$

# Data Model

- Atom - simple atomic value (ie: number or string)
- Tuple – sequence of fields; each field any type
- Bag – Collection of tuples
  - Duplicates possible
  - Tuples in a bag can have different field lengths and field types
- Map – Collection of key-value pairs
  - Key is an atom; value can be any type

$$\left[ \begin{array}{l} \text{'fan of'} \rightarrow \left\{ \begin{array}{l} \text{('lakers')} \\ \text{('iPod')} \end{array} \right\} \\ \text{'age'} \rightarrow 20 \end{array} \right]$$

Key 'fan of' is mapped to a bag containing two tuples, and key 'age' is mapped to an atom 20

# Pig Latin

- FOREACH-GENERATE (per-tuple processing)
  - Iterates over every input tuple in the bag, producing one output each, allowing efficient parallel implementation

```
expanded_queries = FOREACH queries GENERATE
    userId, expandQuery(queryString);
```

- Expressions within the GENERATE clause can take the form of the any of these expressions

$t = \left( 'alice', \left\{ \begin{array}{l} ('lakers', 1) \\ ('iPod', 2) \end{array} \right\}, ['age' \rightarrow 20] \right)$ <p>Let fields of tuple <math>t</math> be called <math>f1, f2, f3</math></p>		
Expression Type	Example	Value for $t$
Constant	'bob'	Independent of $t$
Field by position	$\$0$	'alice'
Field by name	$f3$	'age' $\rightarrow$ 20
Projection	$f2.\$0$	$\left\{ \begin{array}{l} ('lakers') \\ ('iPod') \end{array} \right\}$
Map Lookup	$f3\# 'age'$	20
Function Evaluation	$SUM(f2.\$1)$	$1 + 2 = 3$
Conditional Expression	$f3\# 'age' > 18?$ 'adult': 'minor'	'adult'
Flattening	$FLATTEN(f2)$	'lakers', 1 'iPod', 2

$$t = \left( \text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Let fields of tuple  $t$  be called  $f1$ ,  $f2$ ,  $f3$

Expression Type	Example	Value for $t$
Constant	'bob'	Independent of $t$
Field by position	$\$0$	'alice'
Field by name	$f3$	$[\text{'age'} \rightarrow 20]$
Projection	$f2.\$0$	$\left\{ \begin{array}{l} (\text{'lakers'}) \\ (\text{'iPod'}) \end{array} \right\}$
Map Lookup	$f3\#\text{'age'}$	20
Function Evaluation	$\text{SUM}(f2.\$1)$	$1 + 2 = 3$
Conditional Expression	$f3\#\text{'age'} > 18?$ $\text{'adult'} : \text{'minor'}$	'adult'
Flattening	$\text{FLATTEN}(f2)$	'lakers', 1 'iPod', 2

# Data Model

## User-Defined Functions (UDFs)

Ex: `spam_urls = FILTER urls BY isSpam(url);`

- Can be used in many Pig Latin statements
- Useful for custom processing tasks
- Can use non-atomic values for input and output
- Currently must be written in Java, Python, Javascript



# Flatten

queries:

(userId, queryString, timestamp)

(alice, lakers, 1)  
(bob, iPod, 3)

FOREACH queries GENERATE  
expandQuery(queryString)  
(without flattening)

(alice, {  
 (lakers rumors)  
 (lakers news)  
})  
(bob, {  
 (iPod nano)  
 (iPod shuffle)  
})

(alice, {  
 (lakers rumors)  
 (lakers news)  
})  
(bob, {  
 (iPod nano)  
 (iPod shuffle)  
})

with flattening

(alice, lakers rumors)  
(alice, lakers news)  
(bob, iPod nano)  
(bob, iPod shuffle)

# Load

- Input is assumed to be a bag (sequence of tuples)
- Can specify a deserializer with “USING”
- Can provide a schema with “AS”

```
newBag = LOAD 'filename'  
         <USING functionName() >  
         <AS (fieldName1, fieldName2,...)>;
```

```
Queries = LOAD 'query_log.txt'  
          USING myLoad()  
          AS (userID, queryString, timeStamp)
```

# FOREACH

- Apply some processing to each tuple in a bag
- Each field can be:
  - A fieldname of the bag
  - A constant
  - A simple expression (like  $f1+f2$ )
  - A predefined function (like, SUM, AVG, COUNT, FLATTEN)
  - A UDF

```
newBag =  
    FOREACH bagName  
    GENERATE field1, field2, ...;
```

# FILTER

Select subset of tuples in a bag

```
newBag =  FILTER bagName  
        BY  expression
```

Expression uses simple comparison operators (==, !=, <, >, ...) and Logical connectors (AND, NOT, OR)

```
some_apples =  
    FILTER apples BY colour != 'red';
```

Can use UDFs

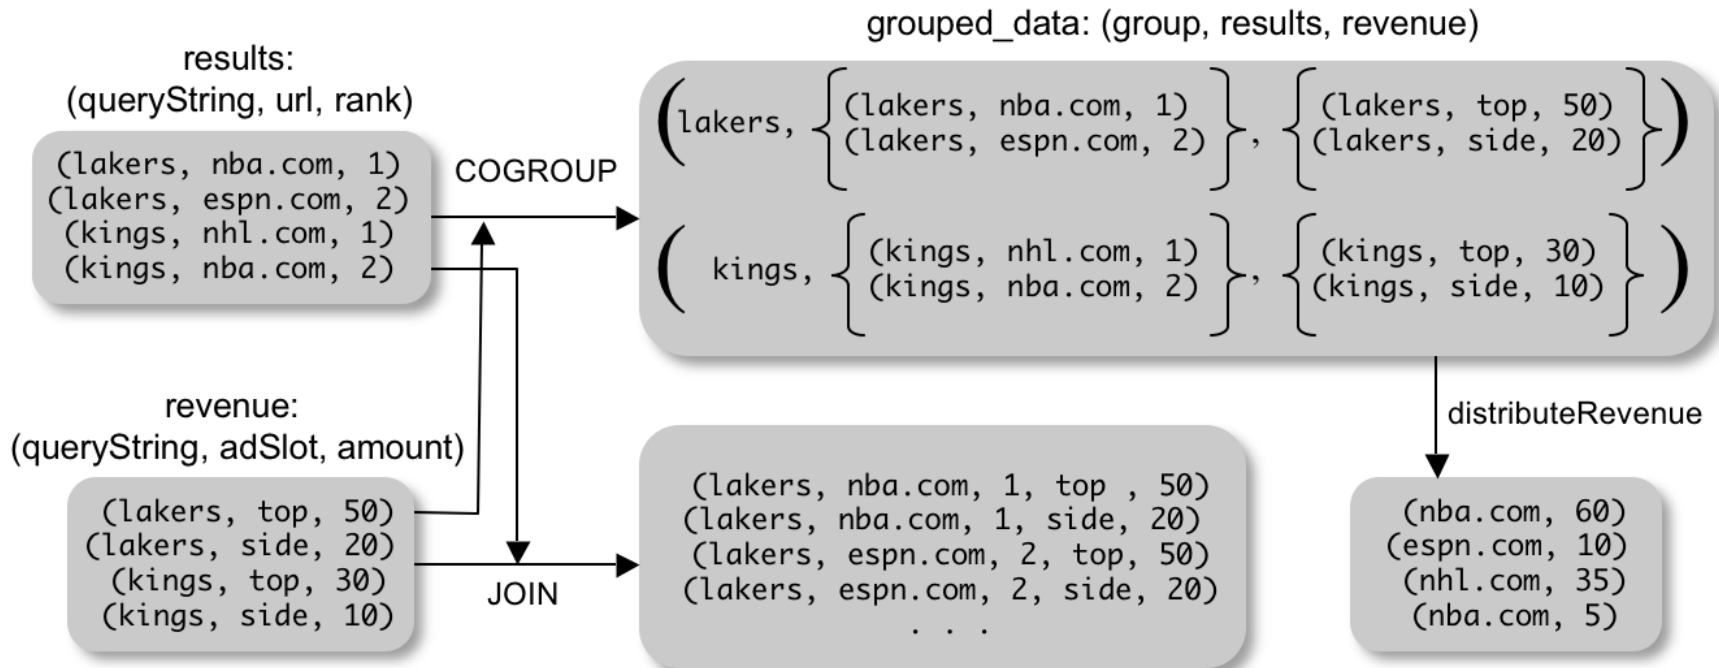
```
some_apples =  
    FILTER apples BY NOT isRed(colour);
```

# COGROUP

Group two datasets together by a common attribute

Groups data into nested bags

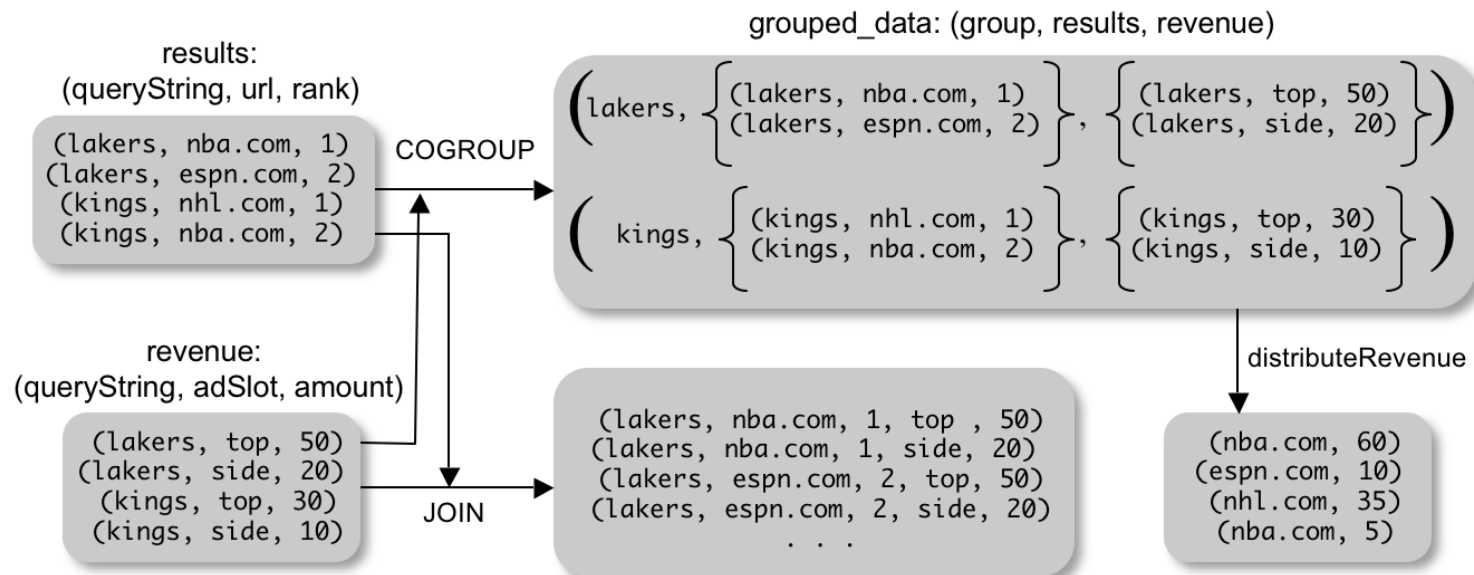
```
grouped_data = COGROUP results BY queryString,  
revenue BY queryString;
```



# COGROUP vs JOIN

```
url_revenues =  
    FOREACH grouped_data GENERATE  
        FLATTEN(distributeRev(results, revenue));
```

- COGROUP takes advantage of nested data structure (combination of GROUP BY and JOIN)
- User can choose to go through with cross-product for a join or perform aggregation on the nested bags



# COGROUP vs JOIN?

To process nested bags of tuples before cross product.

JOIN keyword is still available:

```
JOIN results BY queryString,  
    revenue BY queryString;
```

```
temp = COGROUP results BY queryString,  
        revenue BY queryString;
```

```
join_result = FOREACH temp GENERATE  
                FLATTEN(results), FLATTEN(revenue);
```

# STORE (DUMP)

Output data to a file (or screen)

```
STORE bagName INTO 'filename'  
    <USING deserializer ()>;
```

## Few Other Commands

UNION - return the union of two or more bags

CROSS - take the cross product of two or more bags

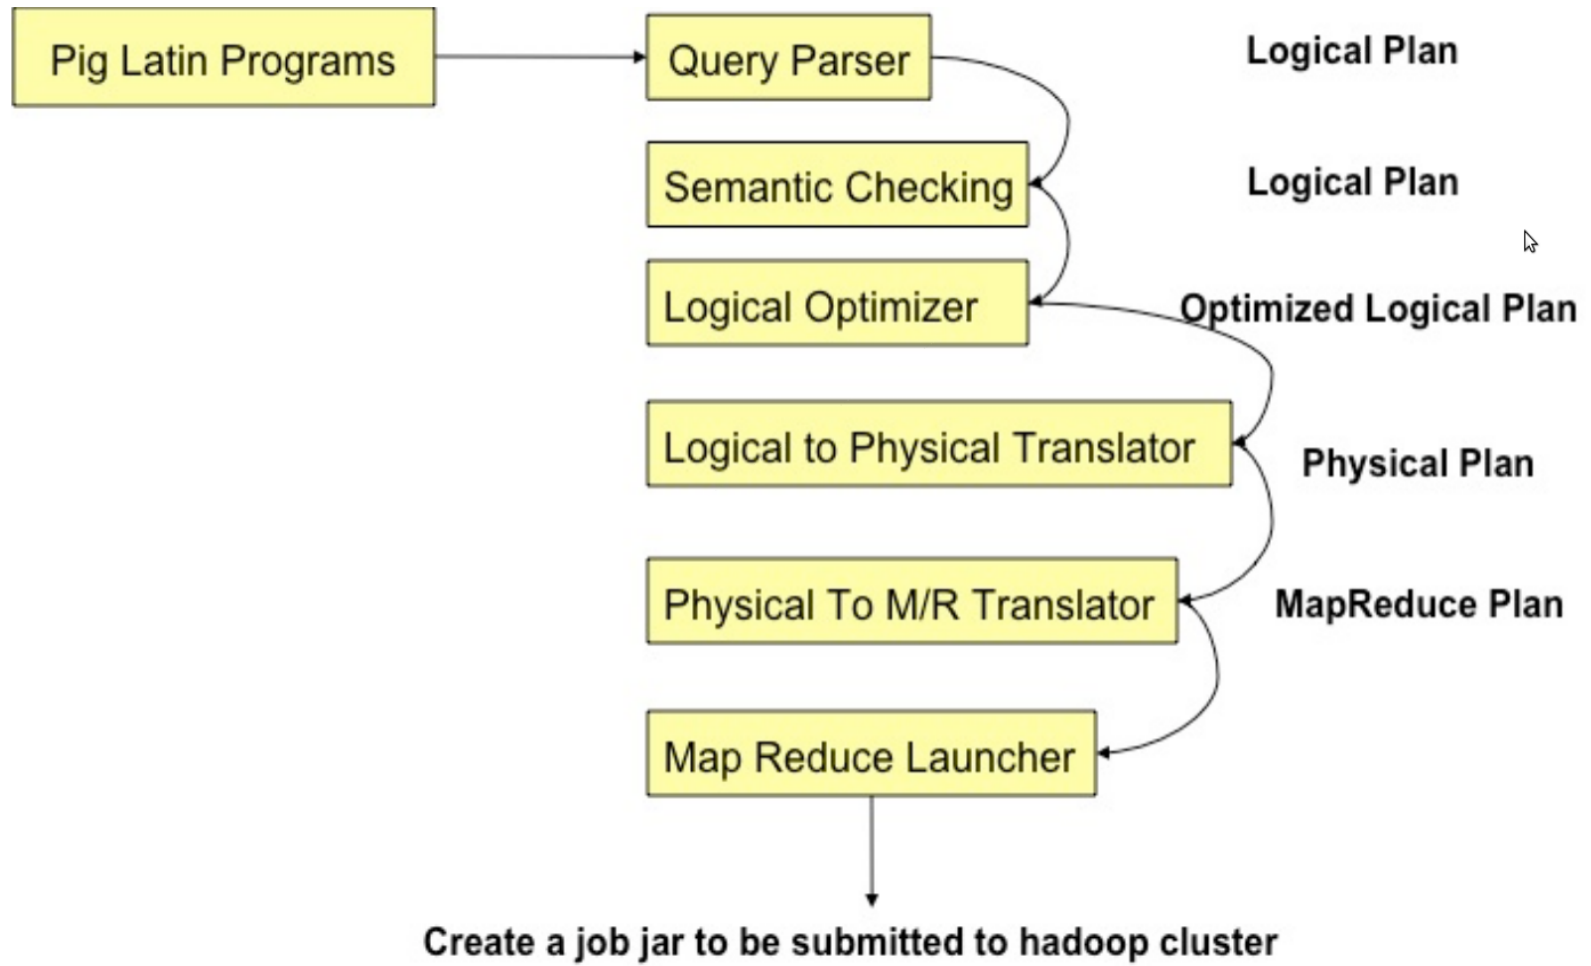
ORDER - order tuples by a specified field(s)

DISTINCT - eliminate duplicate tuples in a bag

LIMIT - Limit results to a subset



# Compilation



# Compilation

Pig system does two primary tasks:



- Builds a Logical Plan from a Pig Latin script
  - Supports execution platform independence
  - No processing of data performed at this stage
- Compiles the Logical Plan to a Physical Plan and Executes
  - Convert the Logical Plan into a series of Map-Reduce statements to be executed by Hadoop Map-Reduce

# Streaming


- Allows data to be pushed through external executables
- Example:  

```
A = LOAD 'data';  
B = STREAM A THROUGH 'stream.pl -n 5';
```
- Due to asynchronous behavior of external executables, each STREAM operator will create two threads for feeding and consuming data from external executables.

# Mapreduce vs Pig



**performance  
1.5x Hadoop**



# Take Home portion of Midterm#1

- Compute the total cost of building a data center that houses 10 ExaBytes of HDD
  - HDDs, networking, racks, building (real-estate), power, cooling, ...
  - You can get the cost of various parts online
  - Assume the datacenter is built in Buffalo
  - State all your assumptions with reference for pricing, etc.
  - Due March 11 at 11:59PM on UBLearns.
  - 25% of your grade for midterm

# Midterm#1

- Parallel/Distributed File Systems
- Statistical Methods
- Some classification
- Mapreduce programming
- R programming
- Pig